

## labset-2

February 28, 2025

*Develop a program to compute the correlation matrix to understand the relationships between pairs of features. Visualize the correlation matrix using a heatmap to know which variables have strong positive/negative correlations. Create a pair plot to visualize pairwise relationship between features. Use california housing dataset.*

```
[27]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[28]: df = pd.read_csv('./housing.csv')
print(df)
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41.0	880.0	129.0	
1	-122.22	37.86	21.0	7099.0	1106.0	
2	-122.24	37.85	52.0	1467.0	190.0	
3	-122.25	37.85	52.0	1274.0	235.0	
4	-122.25	37.85	52.0	1627.0	280.0	
...	...	...	...	...	...	
20635	-121.09	39.48	25.0	1665.0	374.0	
20636	-121.21	39.49	18.0	697.0	150.0	
20637	-121.22	39.43	17.0	2254.0	485.0	
20638	-121.32	39.43	18.0	1860.0	409.0	
20639	-121.24	39.37	16.0	2785.0	616.0	

	population	households	median_income	median_house_value	\
0	322.0	126.0	8.3252	452600.0	
1	2401.0	1138.0	8.3014	358500.0	
2	496.0	177.0	7.2574	352100.0	
3	558.0	219.0	5.6431	341300.0	
4	565.0	259.0	3.8462	342200.0	
...	...	...	...	...	
20635	845.0	330.0	1.5603	78100.0	
20636	356.0	114.0	2.5568	77100.0	
20637	1007.0	433.0	1.7000	92300.0	
20638	741.0	349.0	1.8672	84700.0	
20639	1387.0	530.0	2.3886	89400.0	

```

        ocean_proximity
0          NEAR BAY
1          NEAR BAY
2          NEAR BAY
3          NEAR BAY
4          NEAR BAY
...
20635       INLAND
20636       INLAND
20637       INLAND
20638       INLAND
20639       INLAND

```

[20640 rows x 10 columns]

```
[29]: df.shape #.shape is an attribute of a DataFrame that returns the number of rows
        ↪and columns
```

```
[29]: (20640, 10)
```

```
[30]: df.info() #info() is a method that provides a concise summary of a DataFrame
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
 #   Column                Non-Null Count  Dtype
---  -
0   longitude             20640 non-null  float64
1   latitude              20640 non-null  float64
2   housing_median_age    20640 non-null  float64
3   total_rooms           20640 non-null  float64
4   total_bedrooms        20433 non-null  float64
5   population            20640 non-null  float64
6   households            20640 non-null  float64
7   median_income         20640 non-null  float64
8   median_house_value    20640 non-null  float64
9   ocean_proximity       20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB

```

```
[31]: df.isnull().sum()
```

```

[31]: longitude          0
      latitude          0
      housing_median_age  0
      total_rooms        0
      total_bedrooms    207

```

```

population          0
households           0
median_income        0
median_house_value   0
ocean_proximity      0
dtype: int64

```

```
[32]: df.duplicated().sum()#duplicated().sum() is used to count the number of
      ↪duplicate rows in a DataFrame.
```

```
[32]: 0
```

```
[33]: df.nunique()#nunique() is a method used to count the number of unique values in
      ↪each column of a DataFrame or a specific column.
```

```
[33]: longitude          844
      latitude           862
      housing_median_age    52
      total_rooms          5926
      total_bedrooms       1923
      population          3888
      households           1815
      median_income       12928
      median_house_value   3842
      ocean_proximity       5
      dtype: int64

```

```
[34]: df.fillna({'total_bedrooms':0},inplace = True)#fills missing values in the
      ↪"total_bedrooms" column with 0 and updates the DataFrame in place.
```

```
[35]: df.isnull().sum()#.isnull().sum() is a powerful method used to identify missing
      ↪values in a DataFrame
```

```
[35]: longitude          0
      latitude           0
      housing_median_age    0
      total_rooms          0
      total_bedrooms       0
      population          0
      households           0
      median_income        0
      median_house_value   0
      ocean_proximity      0
      dtype: int64

```

```
[36]: numerical = df.select_dtypes(include = [np.number]).columns
      print(numerical)
```

```
Index(['longitude', 'latitude', 'housing_median_age', 'total_rooms',
      'total_bedrooms', 'population', 'households', 'median_income',
      'median_house_value'],
      dtype='object')
```

```
[37]: correlation_mat = df[numerical].corr()
      print(correlation_mat)
```

	longitude	latitude	housing_median_age	total_rooms	\
longitude	1.000000	-0.924664	-0.108197	0.044568	
latitude	-0.924664	1.000000	0.011173	-0.036100	
housing_median_age	-0.108197	0.011173	1.000000	-0.361262	
total_rooms	0.044568	-0.036100	-0.361262	1.000000	
total_bedrooms	0.068082	-0.065318	-0.317063	0.920196	
population	0.099773	-0.108785	-0.296244	0.857126	
households	0.055310	-0.071035	-0.302916	0.918484	
median_income	-0.015176	-0.079809	-0.119034	0.198050	
median_house_value	-0.045967	-0.144160	0.105623	0.134153	


  

	total_bedrooms	population	households	median_income	\
longitude	0.068082	0.099773	0.055310	-0.015176	
latitude	-0.065318	-0.108785	-0.071035	-0.079809	
housing_median_age	-0.317063	-0.296244	-0.302916	-0.119034	
total_rooms	0.920196	0.857126	0.918484	0.198050	
total_bedrooms	1.000000	0.866266	0.966507	-0.007295	
population	0.866266	1.000000	0.907222	0.004834	
households	0.966507	0.907222	1.000000	0.013033	
median_income	-0.007295	0.004834	0.013033	1.000000	
median_house_value	0.049148	-0.024650	0.065843	0.688075	

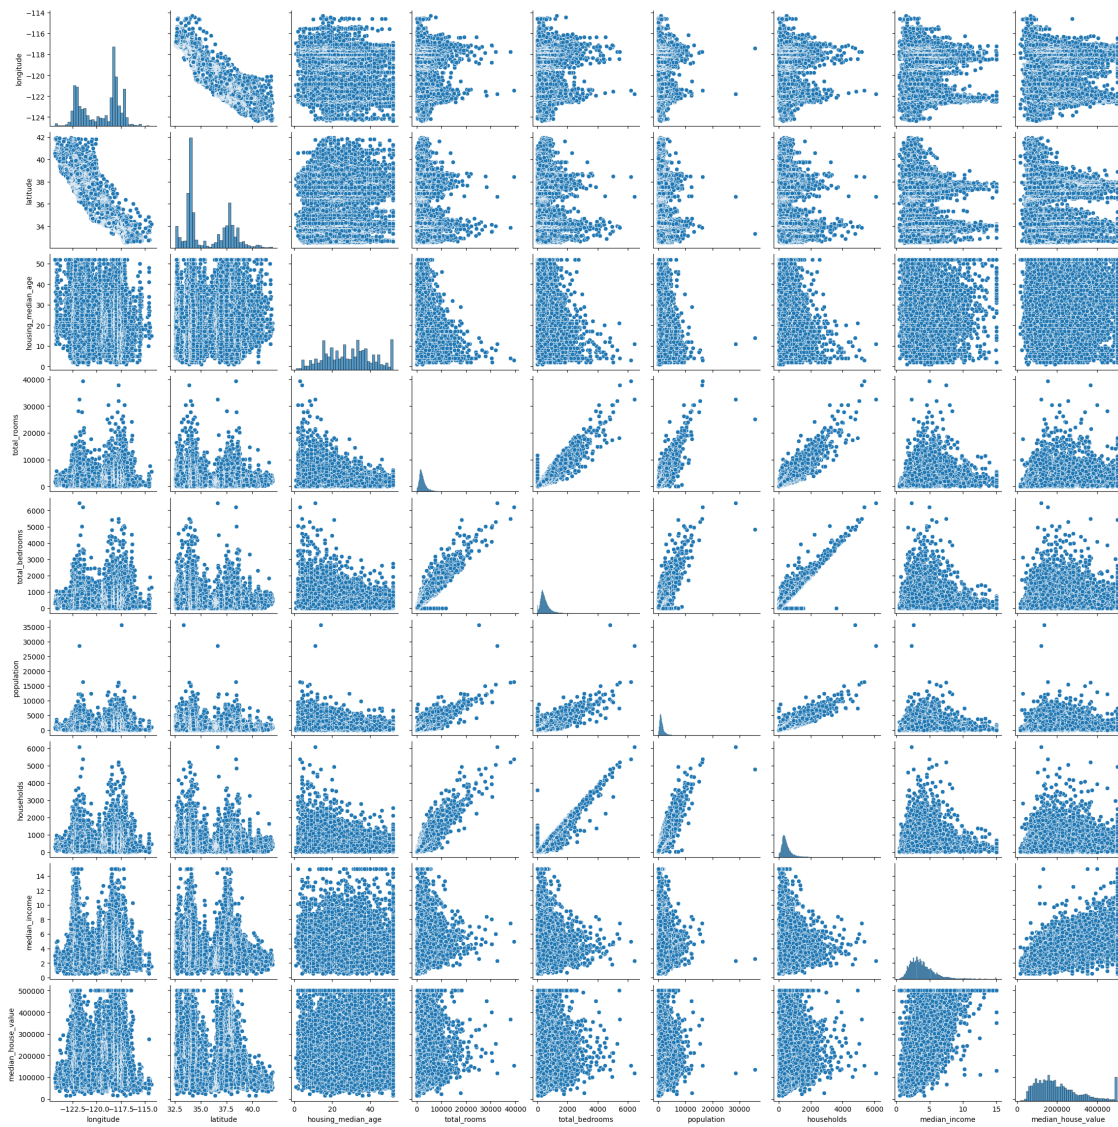
  

	median_house_value
longitude	-0.045967
latitude	-0.144160
housing_median_age	0.105623
total_rooms	0.134153
total_bedrooms	0.049148
population	-0.024650
households	0.065843
median_income	0.688075
median_house_value	1.000000

```
[38]: plt.figure(figsize = (15,10))
      sns.heatmap(correlation_mat,cmap = 'coolwarm',fmt = '.2f',annot = True)
      plt.show()
```



```
[40]: sns.pairplot(df[numerical],diag_kind='hist')
plt.show()
```



[ ]:

[ ]: