

## pol-regress-fl

March 30, 2025

*Develop a program to demonstrate the working of linear regression and polynomial regression. Use Boston housing dataset for linear regression and auto MPG dataset for Polynomial regression*

*Polynomial Regression using Auto MPG Dataset*

```
[33]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import PolynomialFeatures, StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

import warnings
warnings.filterwarnings('ignore')
```

```
[34]: data = sns.load_dataset('mpg')

data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   mpg             398 non-null   float64
 1   cylinders        398 non-null   int64   
 2   displacement     398 non-null   float64
 3   horsepower       392 non-null   float64
 4   weight           398 non-null   int64   
 5   acceleration     398 non-null   float64
 6   model_year       398 non-null   int64   
 7   origin           398 non-null   object  
 8   name             398 non-null   object  
dtypes: float64(4), int64(3), object(2)
memory usage: 28.1+ KB
```

```
[35]: data.nunique()
```

```
[35]: mpg          129
      cylinders     5
      displacement  82
      horsepower    93
      weight       351
      acceleration  95
      model_year    13
      origin        3
      name         305
      dtype: int64
```

```
[36]: data.isnull().sum()
```

```
[36]: mpg          0
      cylinders     0
      displacement  0
      horsepower    6
      weight       0
      acceleration  0
      model_year    0
      origin        0
      name         0
      dtype: int64
```

```
[37]: data['horsepower'].fillna(data['horsepower'].median(), inplace=True)
```

```
[38]: data.isnull().sum()
```

```
[38]: mpg          0
      cylinders     0
      displacement  0
      horsepower    0
      weight       0
      acceleration  0
      model_year    0
      origin        0
      name         0
      dtype: int64
```

```
[39]: # Define X (Feature) and y (Target)
      X = data[["horsepower"]]
      y = data["mpg"]
```

```
[40]: # Split dataset into training and testing sets
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪ random_state=42)
```

```
[41]: # Create polynomial features
degree = 2 # Change the degree of the polynomial
poly = PolynomialFeatures(degree)
X_poly_train = poly.fit_transform(X_train)
```

```
[42]: # Fit a polynomial regression model
model = LinearRegression()
model.fit(X_poly_train, y_train)
```

```
[42]: LinearRegression()
```

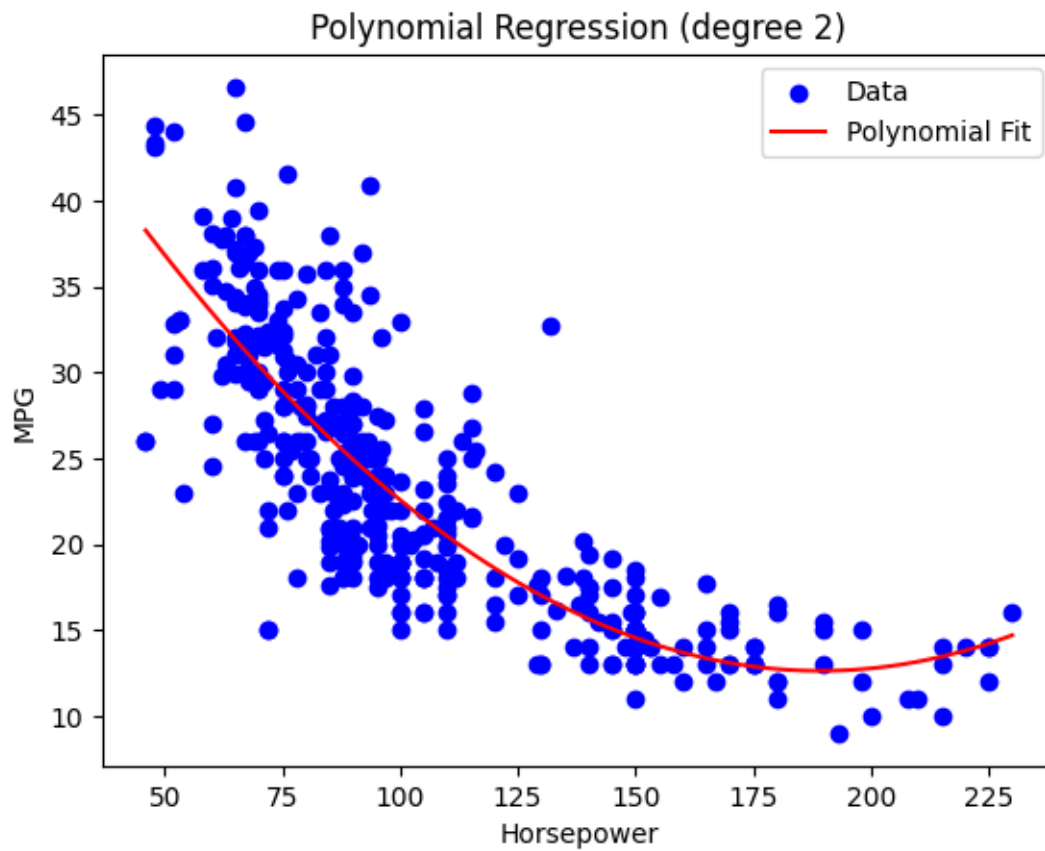
```
[43]: # Make predictions
X_poly_test = poly.transform(X_test)
y_pred = model.predict(X_poly_test)
```

```
[44]: # Calculate performance metrics
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

# Print metrics
print(f'Mean Squared Error: {mse}')
print(f'Root Mean Squared Error: {rmse}')
print(f'R-squared: {r2}')
```

```
Mean Squared Error: 13.941158940364115
Root Mean Squared Error: 3.7337861401483767
R-squared: 0.7407089260880471
```

```
[45]: # Visualize the results
plt.scatter(X, y, color='blue', label='Data')
X_range = np.linspace(X.min(), X.max(), 100).reshape(-1, 1)
X_range_poly = poly.transform(X_range)
y_range_pred = model.predict(X_range_poly)
plt.plot(X_range, y_range_pred, color='red', label='Polynomial Fit')
plt.xlabel('Horsepower')
plt.ylabel('MPG')
plt.legend()
plt.title(f'Polynomial Regression (degree {degree})')
plt.show()
```



[ ]: